Matrix completion: batch size and sampled observations

Each point counts actual sampled training entries, including repeats · Five fresh seeds per batch

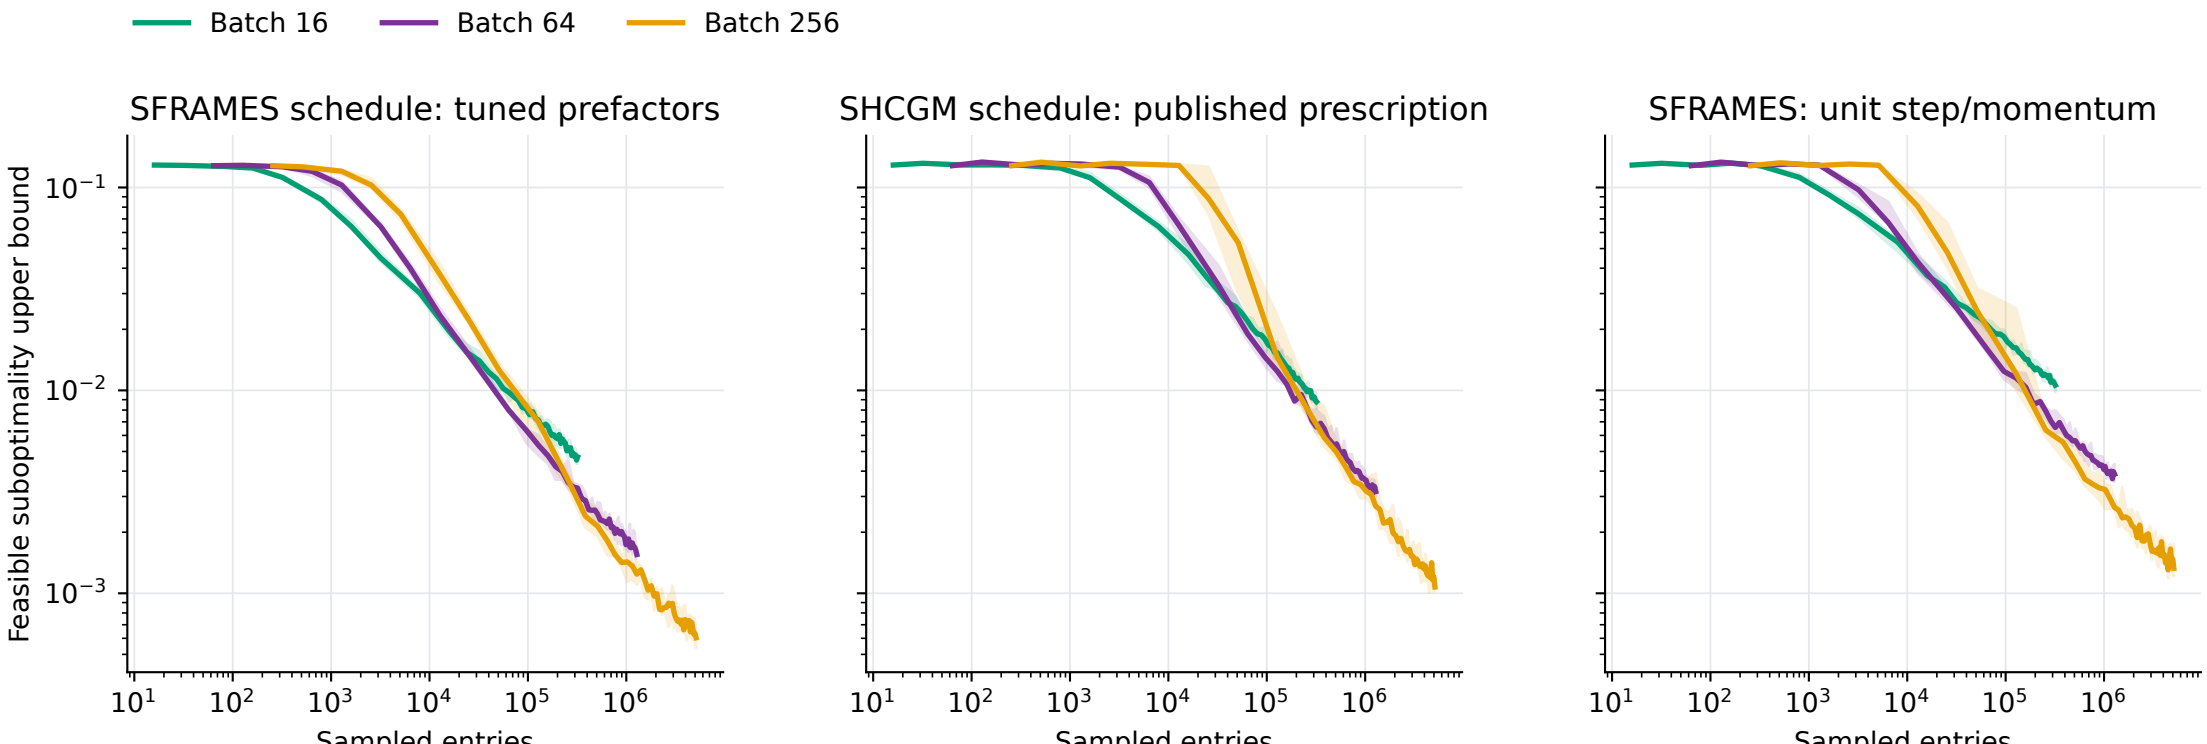

All runs end at 20,000 updates: batch 16 = 320,000 draws; batch 64 = 1,280,000; batch 256 = 5,120,000.

Prefactors were selected at batch 64 and reused unchanged at 16 and 256. Curves stop at their actual observation budgets.

Shaded regions are observed seed ranges, not confidence intervals. Loss is evaluated after radial feasibility repair, used only for diagnostics.