Median feasible GW loss over iterations

Three final seeds per method and initialization · Lower loss is better

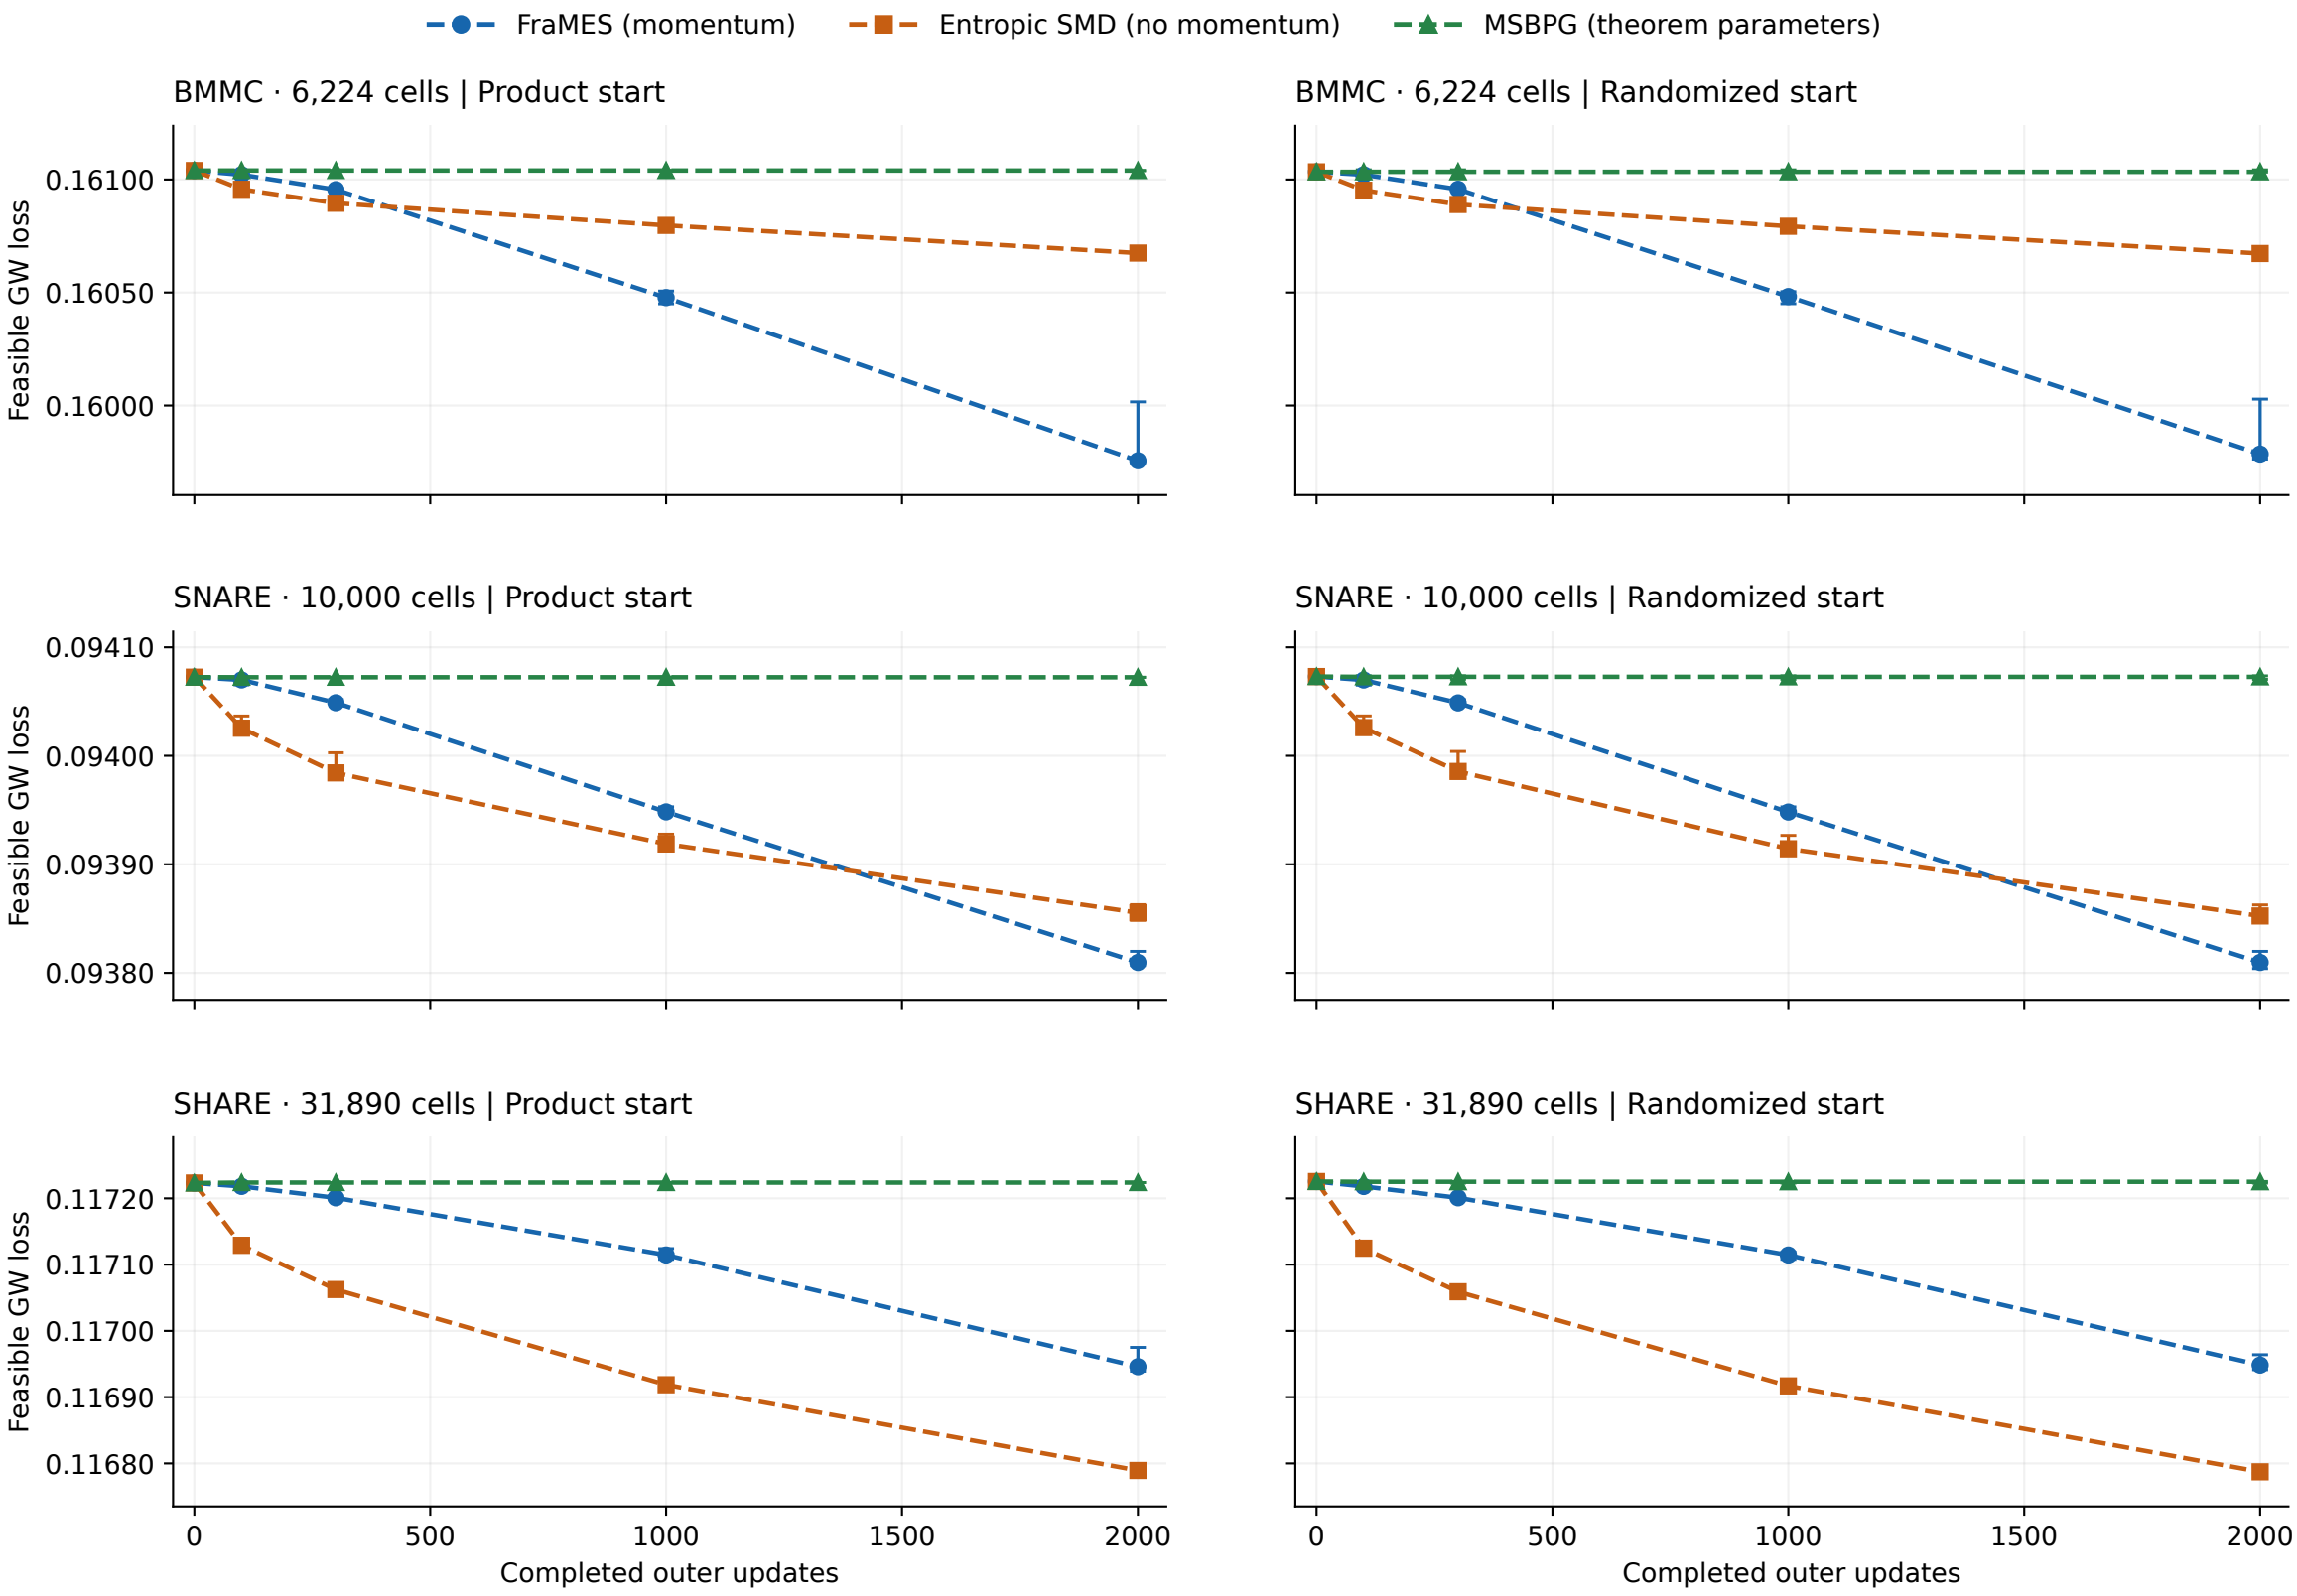

Markers: medians at saved updates 0, 100, 300, 1,000 and 2,000. Bars: interquartile range (n = 3).

Dashed lines connect observed medians as visual guides; no unsaved losses are estimated. Y-scales differ by dataset.

Current rounded outputs; intermediate loss diagnostics use FP32, final losses use FP64. GW loss excludes the KL term.