Quiroga–FRAMES: running on Ruche

5 October 2026, 19:10 CEST | Execution handoff

Ruche execution is underway: smoke and pilot passed execution checks, and the independent n= calibration array 2014199 is running.

Source: exact clean commit 0586c15 is staged in an isolated Ruche checkout. The reviewed setup and source-bundle checksums match. Previous source trees and experiment results are preserved; no GitHub publication was needed.

Completed: Slurm smoke 2014184 (25 s) and calibration-budget pilot 2014189 (22 s), both exit 0. Both pilot methods completed all 256,000 logical row accesses: 2,000 Quiroga updates and 8,000 FRAMES updates. Artifact hashes, JSON/NPZ consistency, source provenance, row budgets and endpoints passed checks.

Pilot physical metrics — one unselected anchor configuration

Quiroga: fidelity 0.1346; full-design RMSE 0.06330; normalized TP violation 0.43866. Recovery gate: not passed.

FRAMES: fidelity 0.7927; full-design RMSE 0.01420; normalized TP violation 0.00683. Recovery gate: not passed.

These are execution/stability checks, not a best-versus-best result. The gate requires fidelity >= 0.99, RMSE <= 1e-6 and normalized TP <= 0.001. The printed adaptive rule retains a calibrated step constant; the paper does not specify every constant needed for exact reproduction.

Calibration array 2014199 is running: one task active, seven queued behind the site's per-user job limit. The first three trials completed and validated. The array allows at most four concurrent tasks, but Ruche currently admits one.

Calibration uses eight channel instances, one shared initialization per channel and three FRAMES sampling streams. A 44-pair-per-channel inventory covers the independent grids: 352 paired executions, with repeated anchors deduplicated for scientific analysis. There are 30 Quiroga candidates and five FRAMES candidates; this is not equal hyperparameter-search effort.

Quiroga tunes a fixed TP penalty and step constant. FRAMES tunes its own initial TP penalty with exponent 0.25, increasing throughout optimization. Both share the physical data, initialization, row budget and recovery gate. Method-native penalized objectives are not compared across methods.

Run and validate the independent calibration grid, deduplicate repeated native-method outcomes, then write the frozen selection manifest before confirmatory validation. The n=2 subset and n=4 phases have not started.

Evidence: archived Slurm logs, pair.json/pair.npz, run_record.json and calibration plan.

No selected validation or held-out comparison is claimed in this report.