Quiroga–FRAMES: Ruche run handoff

5 October 2026, 19:34 CEST

The first corrected calibration channel finished all 44 paired trials. The remaining-channel array is progressing on Ruche.

Completed checks: Ruche smoke 2014184 and the 2,000-row-epoch pilot 2014189 finished successfully. Both methods completed the pilot's 256,000 nominal row accesses. JSON/NPZ consistency, checksums, source provenance and physical metrics recomputed from the saved factors passed.

The untuned pilot did not pass the recovery gate. Quiroga: fidelity 0.1346, RMSE 0.06330, normalized TP 0.43866. FRAMES: fidelity 0.7927, RMSE 0.01420, normalized TP 0.00683. These anchors do not establish a best-versus-best result.

What interrupted the first grid

Array 2014199 exposed a recording bug when an aggressive Quiroga configuration overflowed: final-metric evaluation raised before the failed pair could be archived. Six channel tasks stopped; the remaining two were cancelled. All outputs and error evidence are preserved.

Commit c7a8c97 fixes only failure recording: unavailable failed-endpoint metrics are null with reasons, actual factors and failure status are retained, and valid FRAMES results survive. Optimizer equations, penalties, budgets, gates and the candidate grid are unchanged. A stale launch-path suffix was also corrected before computation.

Validation: 32 focused tests passed; full repository suite 537 passed, 55 skipped. Independent source and launcher reviews approved. The completed Ruche overflow archive and its backup hashes also passed validation on Elite.

Channel 0 (job 2014220) archived all 44 pairs: 43 completed normally and one numerical failure was retained. Its exit 2 deliberately reports that method failure, not an archive error. Array 2014222 is now running channel 1; channels 2–7 are queued.

The grid is unchanged: 30 Quiroga configurations tune a fixed TP penalty and adaptive-step scale; five FRAMES configurations tune their own initial TP penalty, increasing with exponent 0.25. Shared sensing, initialization, row budget and physical gate remain fixed.

Eight channels, one initialization per channel and three FRAMES sampling streams give 352 paired executions. Repeated method-specific results must be deduplicated: 240 unique Quiroga outcomes and 120 FRAMES outcomes. The unequal grid sizes do not imply equal hyperparameter-search effort.

Ruche's cpu_short QoS permits one running job per user. Pending channel tasks therefore wait their turn; they are not evidence of an indefinite resource request.

Finish and validate calibration, deduplicate repeated method-specific outcomes, then freeze the independently selected configurations before confirmatory validation. The n=2 subset and n=4 stages remain next; no final winner is established.

Evidence and exact hashes: run_record.json, launch plans, archived Slurm logs and pair artifacts.
No selected validation or held-out result is claimed. Earlier experiments remain preserved.